

Stanford | Internet Observatory
Cyber Policy Center

Topologies and Tribulations of Gettr

A month in the life of a new alt-network

David Thiel and Miles McCain
Stanford Internet Observatory
August 12, 2021



Contents

1 Whence Gettr?	2
1.1 Design and purpose	3
1.2 Public security issues	3
2 The Gettr API	4
3 Findings: User community	7
3.1 Growth and engagement	7
3.2 Emoji community indicators	8
3.3 Account locations and geographies	11
3.4 New user registrations over time	12
3.5 User bio n-grams	15
3.6 Most popular accounts	17
3.7 Influencers	18
3.8 Top websites, domains, and spam profile prevalence	20
4 Findings: Posts and comments	21
4.1 Predominant post languages	21
4.2 Emoji in posts	23
4.3 Most frequently domains among posted links	24
4.4 Post frequency	25
4.5 Comments	26
5 Sensitive and illegal content	29
6 Conclusions	32

1 Whence Gettr?

Gettr is a far-right social network largely modeled after Twitter. It was launched on July 1, 2021 by former Donald Trump spokesman Jason Miller and promoted and partially financed by Miles Guo,¹ a fugitive Chinese billionaire allied with Steve Bannon² and frequent purveyor of anti-CCP content and COVID-19 misinformation.³ Gettr went through several iterations before its July launch: the platform was under development by a Guo-linked development team since at least early February⁴ under the name “Getter”⁵ and previously launched as the Chinese app “Getome” in June.^{6,7}

Very little is known about Gettr’s moderation practices and Trust and Safety team (if one exists). Its Terms of Service forbid content which is “unlawful, offensive, obscene, lewd, lascivious, filthy, pornographic, violent, threatening, harassing, or abusive, or that violate any right of any third-party.”⁸ In practice, Gettr appears to have few meaningful mechanisms to enforce this. Immediately after launch, SIO researchers observed comment threads bombed with explicit anime, which would have been detectable with sensitive content detection models. These comments were eventually removed, but accounts posting them were not suspended.

Gettr markets itself as a “non-bias social network,” where users are able to post content more freely than Twitter. In some ways, however, Gettr is more restrictive with content policy than the platforms it contrasts itself with—for example, in the app (but not the website), users were initially prevented from posting expletives, presented with an error message of “Sorry, we don’t allow abusive words.” This is a client-side restriction, rather than a server-side filter, and as such was easy for users to bypass. Miller has also suggested that “left of center” users attempting to join the platform were being “caught” and having content removed.⁹ Mechanisms are also in place to allow administrators to add and remove trending topics, suggesting that the company wants to preserve the ability to control narratives itself.¹⁰

Miller claims that Gettr was able to attract 1 million users in the first 3 days after launch, with 1.4 million in the first week.¹¹ Further stories¹² promoted claims that Gettr surpassed 1.5 million users after 11 days; however, according to our analysis, it did not appear to reach this number until the first week of August. The @support account—which every user automatically follows upon creation—shows

¹<https://www.wsj.com/articles/ex-trump-adviser-jason-miller-new-social-app-gettr-backed-by-guo-wengui-tied-foundation-11625199130>

²<https://www.axios.com/steve-bannon-contract-chinese-billionaire-guo-media-fa6bc244-6d7a-4a53-9f03-1296d4fae5aa.html>

³<https://graphika.com/reports/ants-in-a-web/>

⁴<https://archive.is/Yg99Z>

⁵<https://www.politico.com/news/2021/07/01/maga-app-bannon-chinese-billionaire-497767>

⁶<https://gnews.org/1319773/>

⁷<https://twitter.com/emlas/status/1410653796032778249>

⁸<https://gettr.com/terms>

⁹<https://www.thedailybeast.com/jason-miller-sadly-begs-trump-to-join-gettr-says-wed-love-to-have-him>

¹⁰<https://twitter.com/adgodev/status/1410667547335217169>

1.54 million users as of August 9 (see Figure 1).

After the first week, user creation rates have largely plateaued, with growth starting to dip in the latter part of July. Of the 1.5 million users now on Gettr, approximately 372,000 have made a post (including resharing an existing post). Of those, $\approx 45\%$ have made a single post, with $\approx 85\%$ making 10 or fewer. Approximately 234,000 users have made a comment—of those, $\approx 40\%$ have made a single comment, with $\approx 85\%$ making 10 or fewer.



Figure 1: A screen capture of the @support user, showing 1.54 million followers, taken August 9.

1.1 Design and purpose

Gettr was designed to be something of a Twitter “overlay”—essentially, it would import a user’s account from Twitter, but allow extra content that would otherwise violate Twitter’s rules. Twitter follower numbers could initially be imported upon account creation, as a measure of clout; it’s unclear how this is synced going forward and combined with Gettr follower numbers in the UI,¹³ but the API does return Twitter and Gettr followers as separate fields. The user’s Twitter content itself was also initially imported when a user joined Gettr,¹⁴ but this presumably ceased when Twitter cut off Gettr’s API credentials.¹⁵

1.2 Public security issues

From its launch, Gettr was plagued by security and privacy issues. On July 1, security researchers identified several potential security issues in its website,¹⁶ including one that leaked the entire website’s frontend source code.¹⁷ On July 4, a user defaced several prominent accounts on the platform and promoted

¹¹<https://www.foxbusiness.com/politics/jason-miller-gettr-platform-1-million-users-after-launch>

¹²<https://www.washingtonexaminer.com/washington-secrets/gettr-explodes-passing-1-5-million-users-in-just-11-days>

¹³According to reverse-engineered source code, Gettr followers are (or were) added to Twitter followers; however, SIO observed some users’ displayed follower counts decreasing and increasing drastically for no obvious reason, while their actual follower counts returned by the API did not.

¹⁴<https://www.politico.com/news/2021/07/01/gettr-trump-social-media-platform-497606>

¹⁵<https://twitter.com/JasonMillerinDC/status/1413902414411935746>

their own Twitter account.¹⁸ Then, two days later, Gettr experienced its first clear data breach: threat actors were able to use the Gettr API to extract the private information—including email and birth year—from 90,000 users.¹⁹

In a statement about the breach, Jason Miller acknowledged the issue and said that “GETTR takes cybersecurity seriously and has undertaken another round of security testing by a “white hat” security firm to ensure safety.”¹⁹ Despite Miller’s assurance of Gettr’s security and claim that Gettr had already undergone a security audit, the site remains beset with security issues; for example, several users uploaded images hosted at arbitrary external URLs, leaking the IP address and User Agent of any user who views certain posts, comments, and profiles. The Stanford Internet Observatory has identified several public posts possibly intended to extract user IP address information.

2 The Gettr API

We collected our data from Gettr via its publicly accessible API using an open source Gettr data archival tool.²⁰ We collected user data on August 1, 2021, exactly one month after Gettr’s initial public launch, and post/comment data periodically over the course of July 2021.

Records returned by the Gettr API offer some insight into the design and operation of the site itself; however, given that this API is not publicly documented, our interpretation of some of the fields are speculative. Of note, in the objects returned for user profiles, both follower/following counts for Gettr itself and those for a user’s linked Twitter account are presented, and presumably used to calculate a user’s “influence” or adjusted follower count. There is also an `infl` field which appears to designate different users as having different tiers of “influence.” We analyze these influence levels further in Section 3.7 on page 18.

User objects also contain a `meta` field, which currently contains only a boolean field called `auto_connect`. This only appears (set to `False`) on 47 accounts, all created during the first 10 days of Gettr’s operation, and all of which are “influencers.” This may be related to syndication of content from Twitter before Gettr’s Twitter API access was revoked.

¹⁶<https://www.vice.com/en/article/z3xqva/gettr-is-the-trump-teams-buggy-leaky-twitter-clone>

¹⁷<https://github.com/sudofox/rtteg>

¹⁸<https://techcrunch.com/2021/07/06/gettr-trump-social-network-hack-defaced/>

¹⁹<https://www.vice.com/en/article/dyv44m/hackers-scrape-90000-gettr-user-emails-surprising-no-one>

¹⁹Ibid.

²⁰<https://github.com/stanfordio/gogettr>

Key	Value
<code>update</code>	User update date
<code>_t</code>	Always <code>unix</code> (for user information)
<code>_id</code>	Internal, static id string
<code>nickname</code>	Display name
<code>username</code>	Appears always same as <code>_id</code>
<code>username</code>	Username, allows mixed case
<code>status</code>	User “trust” status (e.g. unverified, quick signup, active, suspended)
<code>create_date</code>	User creation date
<code>language</code>	User UI language setting
<code>following</code>	Following count
<code>followers</code>	Follower count
<code>profile_picture</code>	User profile picture
<code>banner_image</code>	Banner/background image
<code>likes</code>	How many posts the user has liked
<code>description</code>	User description/bio
<code>location</code>	Stated user location
<code>website</code>	Arbitrary URL
<code>likes</code>	How many comments the user has liked
<code>shares</code>	Number of posts the user has shared
<code>pin_posts</code>	An array of the user’s pinned posts
<code>infl</code>	Influencer level; null or 1 through 5
<code>twitter_follows</code>	Number of people the user follows on Twitter
<code>twitter_followers</code>	Number of people following the user on Twitter
<code>meta</code>	An array currently boolean <code>auto_connect</code>

Table 1: Structure of a user object. `profile_picture` and `banner_image` are apparently intended to be relative links to an image hosted on Gettr, but can be absolute URLs to an external resource (see Section 1.2 on page 3). Times are in milliseconds since UNIX epoch.

User posts and comments similarly contain both engagement data for the Gettr post itself as well as the Twitter engagement and retweet information if the post was mirrored from Twitter. This is not clear in the user interface, potentially making it seem like more content and interaction is present on Gettr than there actually is.

Key	Value
acl	Whether the post is public or not, does not currently appear to have any effect
_id	Post or comment unique identifier (base-36 string prepended with 'p' for posts or 'c' for comments)
txt	Textual body of the post or comment
htgs	Array of hashtags used
uid	User id of posting or commenting user
cdate	Post creation date
update	Post update date
_t	Type ("post" for posts, "cmt" for comments)
uinf	Array of user metadata for posting user
shrdpst	If the post is a reshare of another post, that post's data
s_pst	If post, id and engagement data for the post
s_cmst	If comment, array containing engagement data for the comment
post	If comment, array containing original post information, sans engagement
puid	If comment, the user id of the post being replied to
vis	Always null or "p"
vid_wid	Width of posted video
vid_hgt	Height of posted video
imgs	Array of relative URLs of posted images
meta	Geometry metadata
main	Image URL
vid_dur	Duration of video
vid	Relative URL to m3u8 of video
ovid	Relative URL to actual video
twt_rtourl	If retweet of a tweet, link to that tweet's URL
prevsrc	URL that the story preview was generated from
twt_rtpst	If retweet, array of metadata of both tweet and user
ipid	Twitter identifier, either post or user
twt_lkbpst	Twitter likes for post
twt_shbpst	Twitter shares for post
is_pinned	Is a pinned post (on Gettr)
utgs	Array of users that were tagged
ttl	If a post of a URL, title of that story/page
dsc	Description/subtitle of above
previmg	Preview image of the post (remote hotlink)
rpstIds	If a repost of a Gettr post, its id
rusrIds	If a repost of a Gettr post, user id of the original poster
pid	Post that a comment pertains to
ddate	Date, almost always null
mtgs	For comments, always null

Table 2: Structure of a post/comment object. In the case of posts imported from Twitter, creation/update dates are of the tweet, not the post on Gettr itself. Several of these fields are null for comments.

3 Findings: User community

3.1 Growth and engagement

Figure 2 below and Figure 3 on the following page show two different views of Gettr new user registrations, comments and posts over time. While new user registrations have declined over time, posts and comments remain fairly static, indicating that Gettr may have reached something of an activity plateau, or that many users are not interacting with the site or are consuming the news feed passively. There is a notable sudden jump in comment activity around July 13, where comments begin to track more closely with posts.

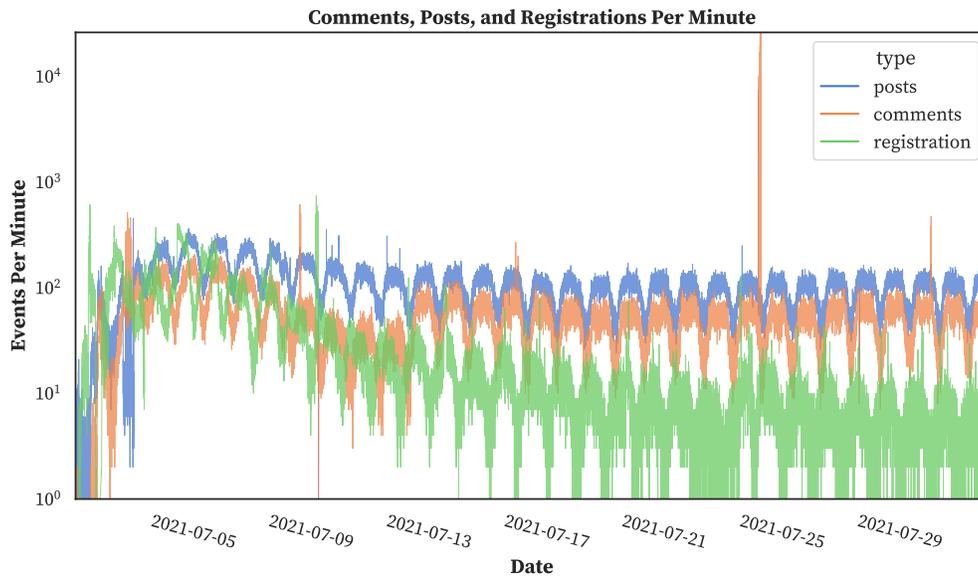


Figure 2: Registrations, posts, and comments per minute on Gettr during the month of July 2021 (log scale).

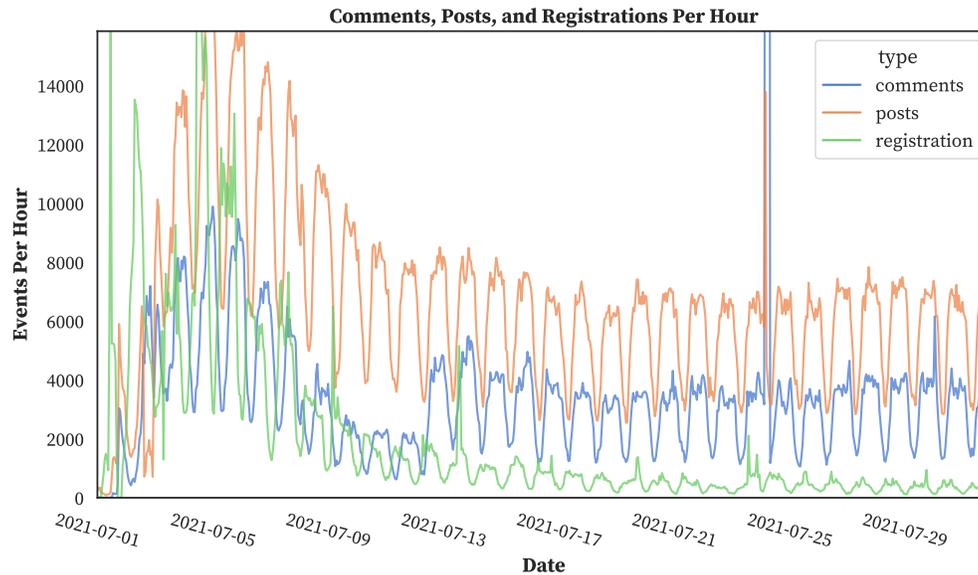


Figure 3: Registrations, posts, and comments per hour on Gettr during the month of July 2021 (linear scale). The y-axis has been capped at $10^{4.2}$ for readability.

3.2 Emoji community indicators

Social media users often use emoji to intentionally signal communities they identify with—for example, by including an American or Brazilian flag in their profile. The Stanford Internet Observatory has previously used the frequency of emoji in public account information to characterize communities within platforms.²¹ We find that 20,650 Gettr users included a United States flag emoji in their profile description, followed by 11,350 users who included a Brazilian flag emoji in their description, indicating significant American and Brazilian communities. The “blue heart” is commonly seen with American pro-police sentiment. More unexpectedly for a right-wing network, the rainbow flag is highly represented, with several hundred users including it in their description.

²¹ <https://cyber.fsi.stanford.edu/io/news/sio-parler-contours>

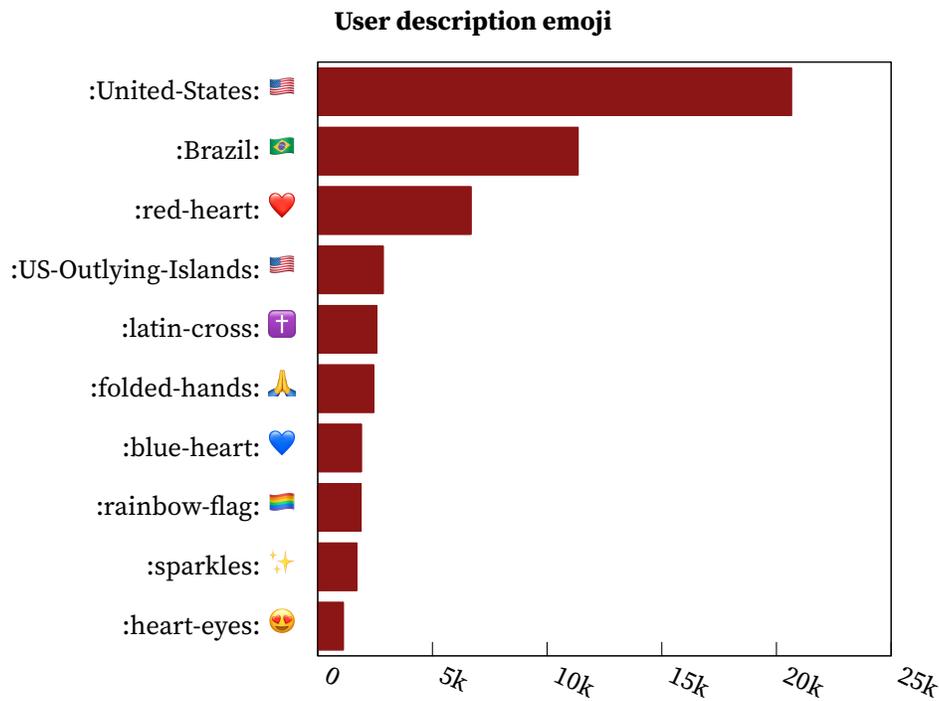


Figure 4: Most frequently used emoji in user bio/description fields.

The most common emoji used in user nicknames (the Gettr equivalent of a display name) are variations of a “check mark” emoji, frequently used to make an account appear “verified,” with varying degrees of success—examples are shown in Figure 6 on the next page. The second most common emoji was the Brazilian flag, followed by the United States flag—this may be explained by the perception that Gettr is a US-centric social network, thus making US identity implicit and the Brazilian flag a more explicit community identifier for certain non-US users. The Saudi Arabian flag (344 users), Israeli flag (325 users), German flag (163), and Japanese flag (129 users) were all also among the top 25 most common emoji in user nicknames.

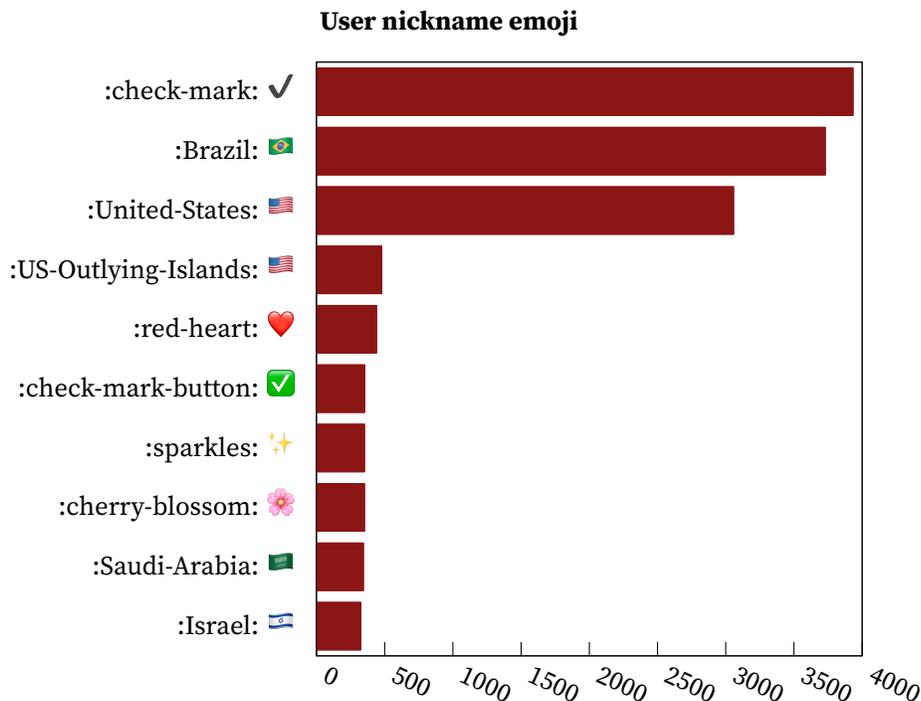


Figure 5: Most frequently used nickname emoji. US Outlying Islands is visually identical to the mainland US flag, and tends to be selected by users of certain Android devices.



Figure 6: A verified account (Naomi Wolf) with Gettr’s red “V” in the nickname, and three accounts attempting to add credibility by using a checkmark emoji in their nickname. The Daily Caller currently does not have an official account on Gettr.

3.3 Account locations and geographies

Gettr allows users to add a location to their profile. Gettr does not verify these locations in any way—it is simply a text field, and some users enter locations such as “definitely not the mothership” and “Swamp Adjacent,” while others enter actual geographic locations such as “Ijuí, Brasil” and “Ontario, CA.” We geoparsed each self-declared location to extract the country and normalize the data, showing that approximately 150,000 users identified as being in the United States, 80,000 in Brazil, 9,000 in Canada, and 7,000 in Germany. A further breakdown can be seen in Figure 7.

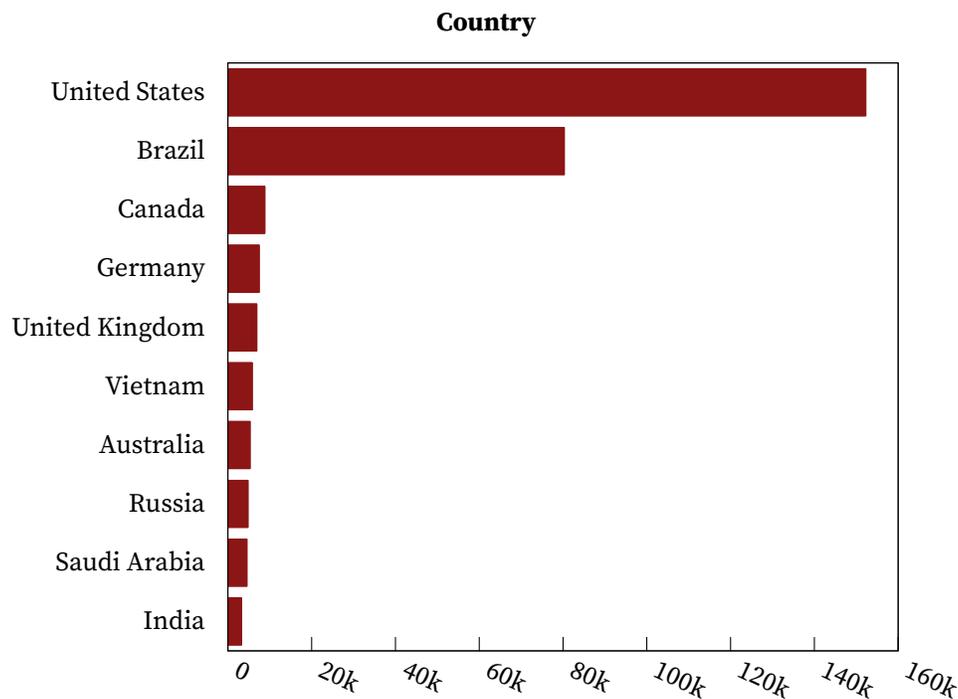


Figure 7: Users by normalized declared location.

Information on language demographics is not obtainable directly, as the `lang` field appears to only represent the display language for the Gettr UI that the user has selected, rather than the primary language of a user’s content. Many users stay with default English, and the range of supported UI language translations is limited. To get a more accurate breakdown of user primary languages, we used the CLD3²² neural network model to detect the languages used in the user “description” field. The results show a significant Portuguese-, Spanish- and Arabic-speaking userbase, similar to previous analyses of right-wing networks such as Parler.²³

²²<https://github.com/google/cld3>

²³<https://cyber.fsi.stanford.edu/io/news/sio-parler-contours>

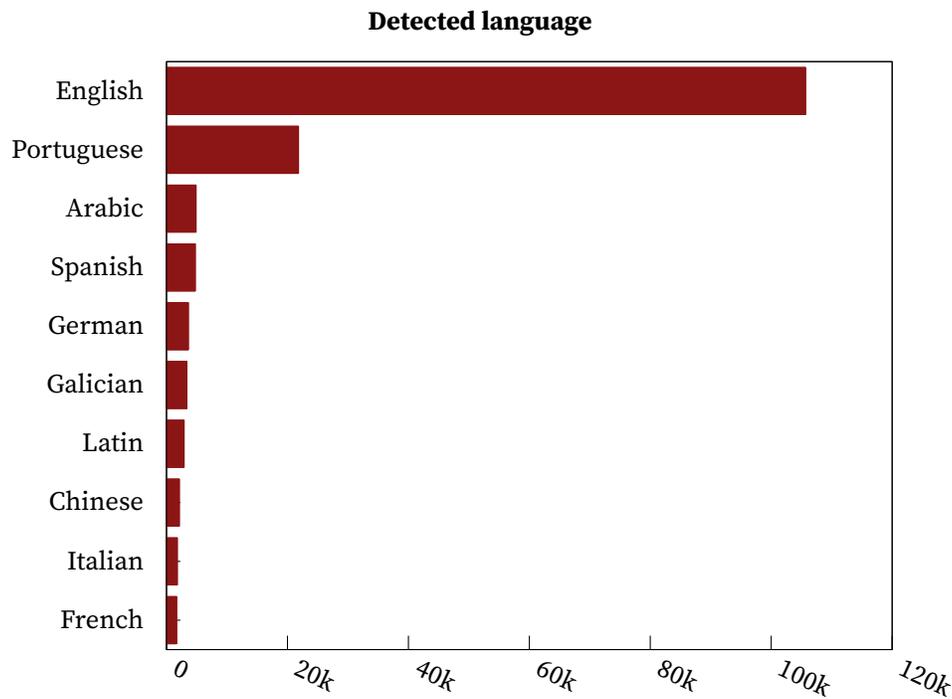


Figure 8: Detected language in user bio, with high-false-positive languages excluded. Note that those detected as Latin or Galician are a mix of Spanish and Portuguese, along with many of those detected as Italian.

3.4 New user registrations over time

Gettr grew quickly in the days after launch, with registrations hitting more than 600 per minute during a peak on July 1, 2021. Growth remained relatively stable until Gettr’s “official launch” on July 4, at which point registrations began to slow. On July 7, Gettr’s support page acknowledged that in-app registrations were broken, perhaps partially explaining the initial decrease in new users.

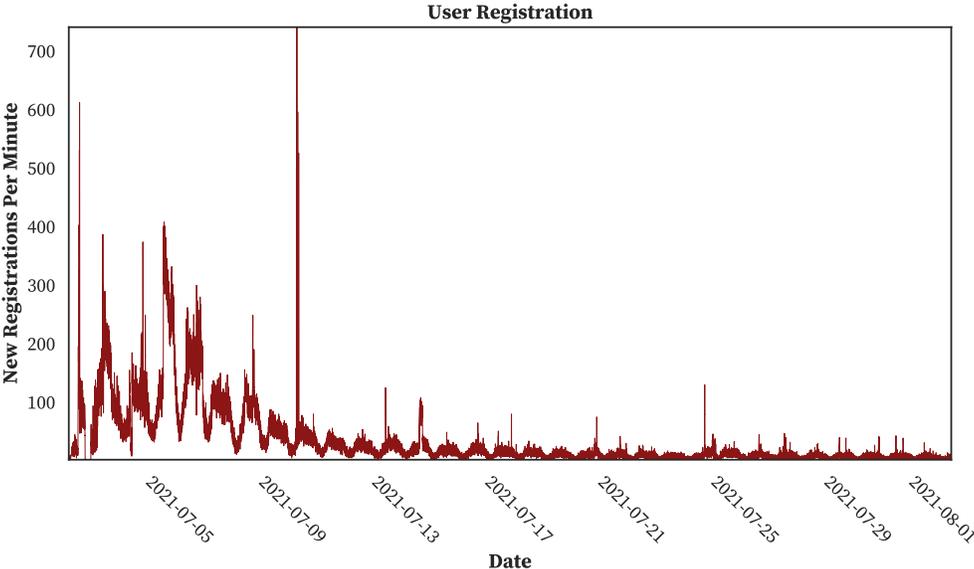


Figure 9: Total new user registrations over time.

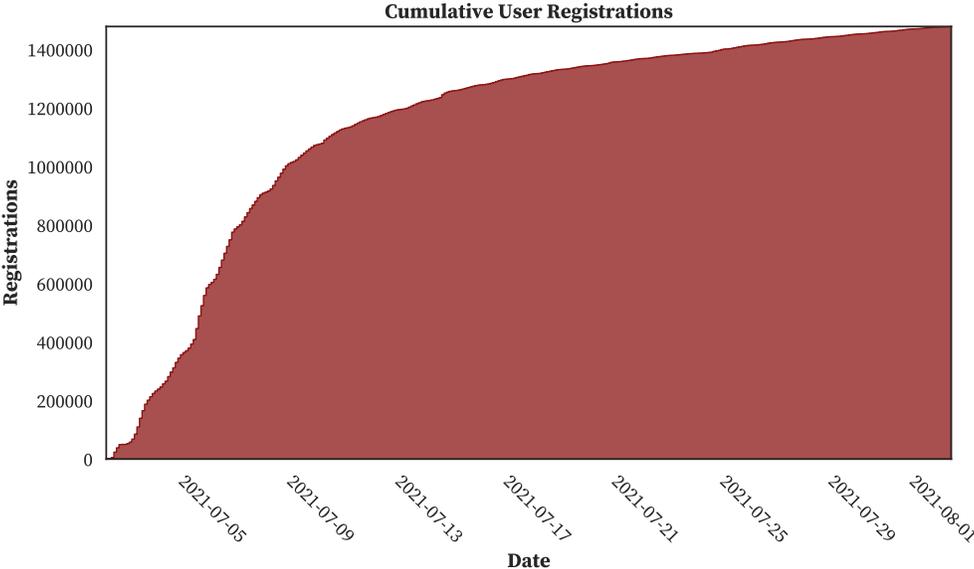


Figure 10: Cumulative user registration over time.

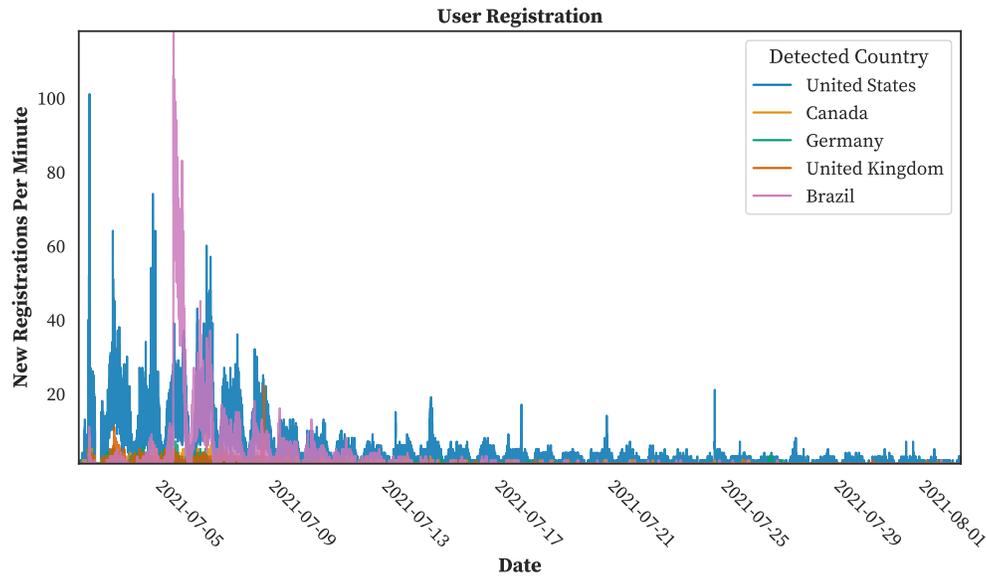


Figure 11: New user registration, broken down by normalized declared country.

Disaggregating new user registrations by geoparsed location reveals that growth over Gettr’s first three days was driven by US-identifying users and, to a smaller extent, British- and German-identifying users. On July 4, hundreds of Brazilian-identifying users joined Gettr, temporarily surpassing the rate of American-identifying user registrations. While the number of Brazilians joining Gettr receded in the following days, Brazilians—and, specifically, Bolsonaro supporters—remain a prominent group on Gettr.

3.4.1 New user registrations by detected language

Disaggregating new user registrations by detected language reveals similar trends. While the vast majority of initial user registrations were English speaking, a smaller but significant amount of initial registrations were Arabic speaking. On July 5, Portuguese-speaking registrations temporarily surpassed English-speaking ones. This is likely due to Flávio Bolsonaro, son of Brazilian president Jair Bolsonaro, announcing his Gettr account on Twitter (see Figure 13 on the next page).

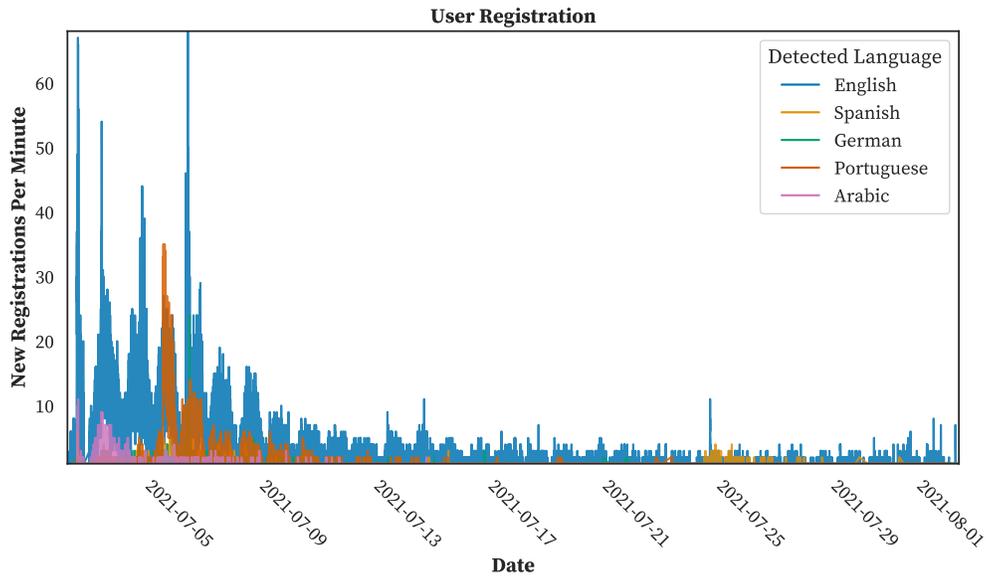


Figure 12: New user registrations by language detected in user description.



Figure 13: Flávio Bolsonaro announcing his new presence on Gettr.

3.5 User bio n-grams

Bigram and trigram analyses are often used to gauge the character and overall sentiment of text. We analyzed user biographies, posts and comments to find the most frequently occurring n-grams on the platform.

Bigrams and trigrams in user descriptions reveal a mix of conservative and right-wing sentiment in both English and Portuguese, along with a significant amount

of self-promoting spam and pornographic accounts.

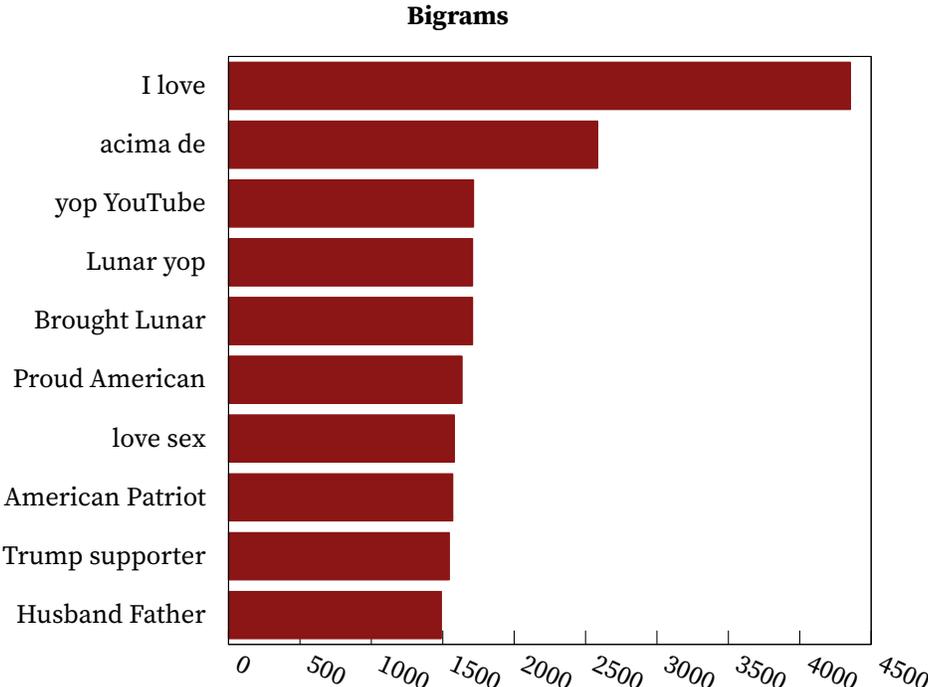


Figure 14: Top bigrams in user descriptions (“bios”).

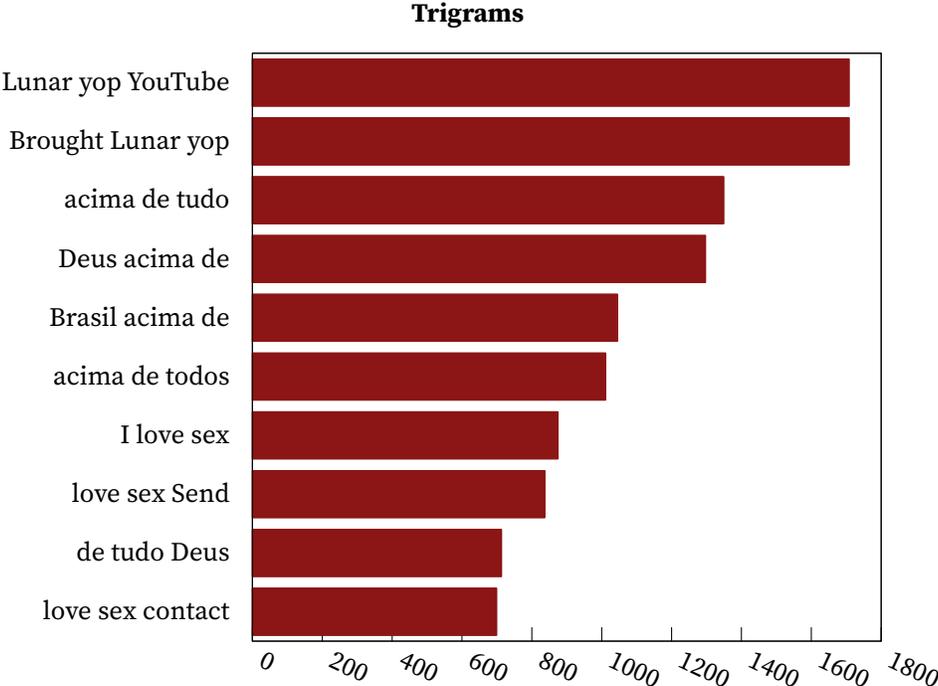


Figure 15: Top trigrams in user descriptions.

3.6 Most popular accounts

The most popular Gettr accounts by number of followers are American conservatives: the most followed account is Mike Pompeo, followed by Newsmax, Dinesh D’Souza, and Steve Bannon’s “War Room.” Other notable popular accounts include “presdonaldtrump” (who impersonated President Trump and was suspended after reaching more than 100,000 followers), Jair Bolsonaro, and Flavio Bolsonaro. While Gettr allows users to “import” their follower counts from Twitter, this analysis only considers each user’s Gettr followers (and excludes followers they “imported” from Twitter).



Figure 16: Mike Pompeo’s profile on Gettr. Despite Gettr only having approximately 1.5M users at the time, Pompeo’s account appears to have 1.61M followers. Gettr is adding Pompeo’s approximately 1.1M Twitter followers to his 500K Gettr followers to compute his overall follower count.

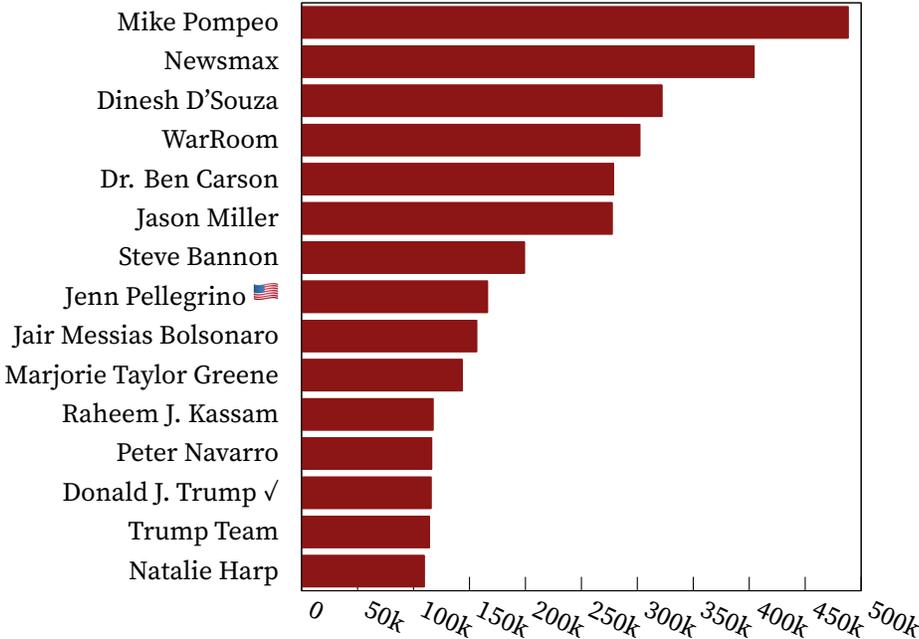


Figure 17: Most followed users (using Gettr metrics, not imported Twitter counts).

3.7 Influencers

As shown in Table 1 on page 5, Gettr users all have an `infl` classification which categorizes them as different types of influencers. For most users, this field is null, but ≈ 1600 users have a value of 1, making them the most common “influencer” type (see Table 18 on the next page). There are only 15 type 4 influencers (Table 19 on the following page) which include investors and those involved in the commercial operation of Gettr itself. There are also two outliers: the `@support` account is type 5, the highest level of influencer (this account is automatically followed by all new users), and “Joe Pags Pagliarulo” / `@joetalkshow`, is type 2.

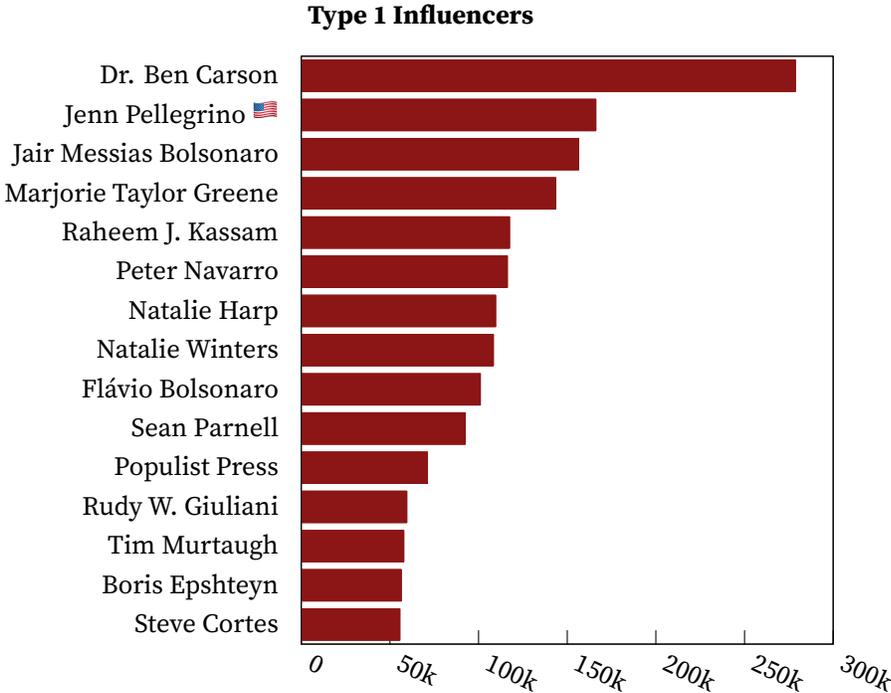


Figure 18: Top type 1 influencers out of 1587, sorted by follower count.

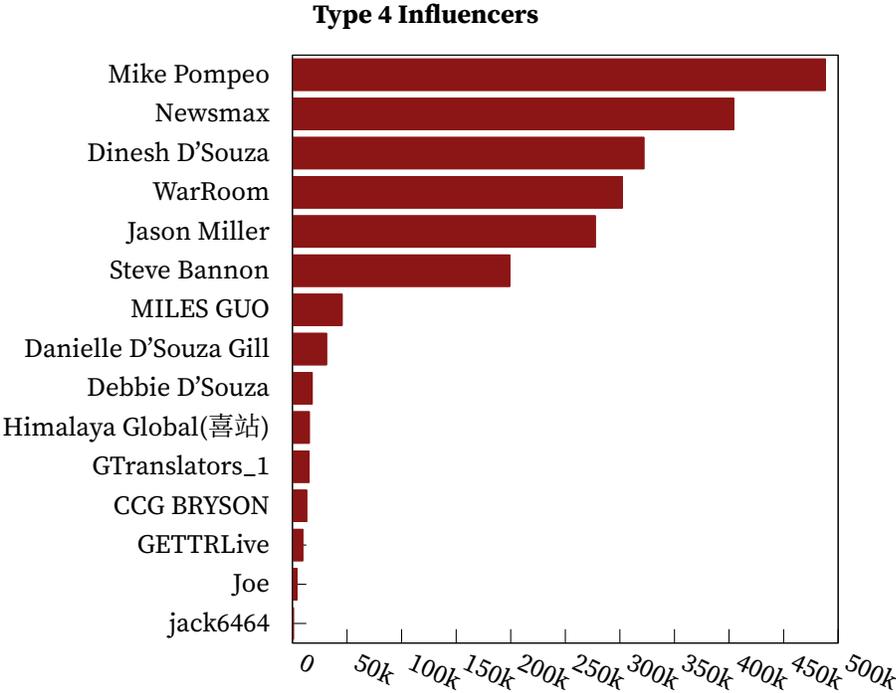


Figure 19: Type 4 influencers, out of 15.

3.8 Top websites, domains, and spam profile prevalence

Gettr allows users to link to a website in their profiles. Of the nearly 1.5 million users on Gettr, less than 5% use this feature. While some users use this feature to link to their website (e.g., Sean Spicer links to SeanSpicer.com) or to signal personal values or political affinity (such as linking to DonaldJTrump.com), other users use this feature for spam and advertisements. Indeed, the most frequent ten URLs are all either spam or troll content. The most common URL is a broken link to YouTube associated with nearly 1,750 profiles that appear to have been created as part of an automated spam campaign (notably, most of the accounts with this link in their profile have been suspended). The second and third most common URLs are adult scam websites. Also in the ten most common URLs is a link to Rick Astley’s “Never Gonna Give You Up” on YouTube (to assist the phenomenon of “rickrolling”). The dominance of spam in the most common URLs indicates a fair number of trolling or otherwise inauthentic accounts.

URL	Count
youtube[.]com/watch?v=LITERALLY-ANYTHING-HERE	1739
bit[.]ly/3dRB7Dw	862
vipdeit[.]com	699
pornhub[.]com	614
turnmeon[.]club	526
fbi[.]gov	345
www.pornhub[.]com	179
sites.google[.]com/view/domenclature/home	121
youtu[.]be/dQw4w9WgXcQ	110
girlsplus.mcdi[.]me	106
gettr[.]com/user/patriotprotect	88
www.donaldjtrump[.]com	87
e621[.]net	82
donaldjtrump[.]com	78
xvideos[.]com	67

Table 3: Most common URLs in the user “website” field. “LITERALLY-ANYTHING-HERE” is the actual text provided by the user, not an indication of a placeholder.

Examining the most common *domains* in user profiles reveals broader trends. Twitter.com is the most frequently linked domain, suggesting that many users are not completely abandoning Twitter. Many link to Instagram—either as “www.instagram.com” or “instagram.com”—and to Telegram groups at the “t.me” domain. Many also link to pornography—based on examination of these profiles, this appears to be predominantly trolling behavior, rather than adult content marketing.

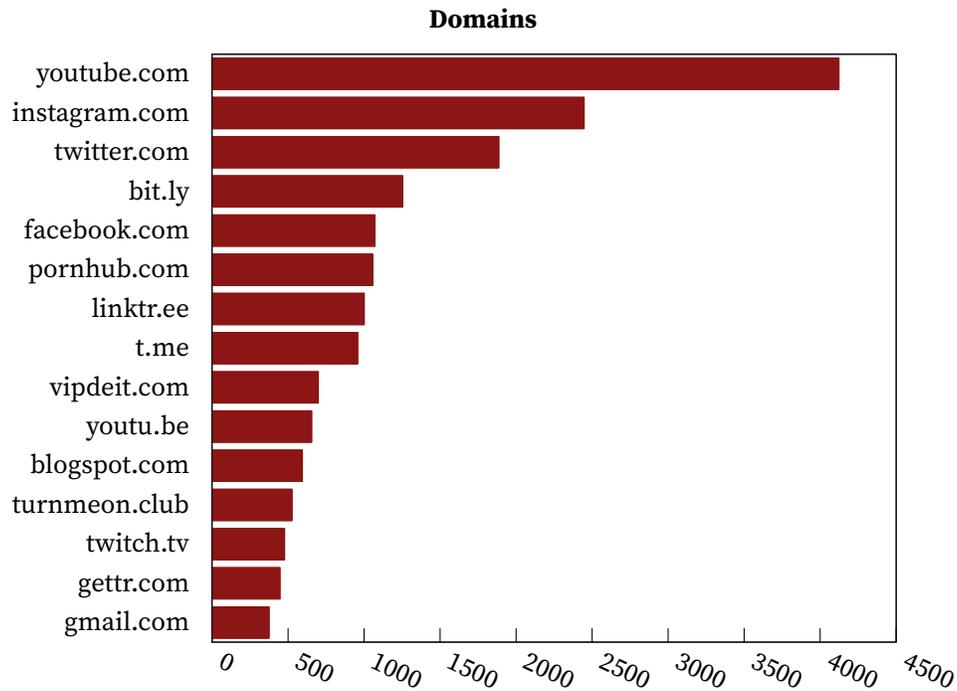


Figure 20: Top domains in user website field.

4 Findings: Posts and comments

4.1 Predominant post languages

Analyzing the languages of user posts shows that the users actually producing content are not necessarily representative of the overall userbase (as detected with other community analysis methods such as declared location, emoji or detected language of user description). For example, despite Arabic speaking users being a smaller part of the user population (as detected by user description language), Arabic posts are almost twice as common as those in Portuguese. There are also a significant amount of Chinese-language posts, considering that only ≈ 2100 users have any Chinese in their description. (The outsized presence of Arabic and Chinese compared to the their userbase holds for comments as well.)

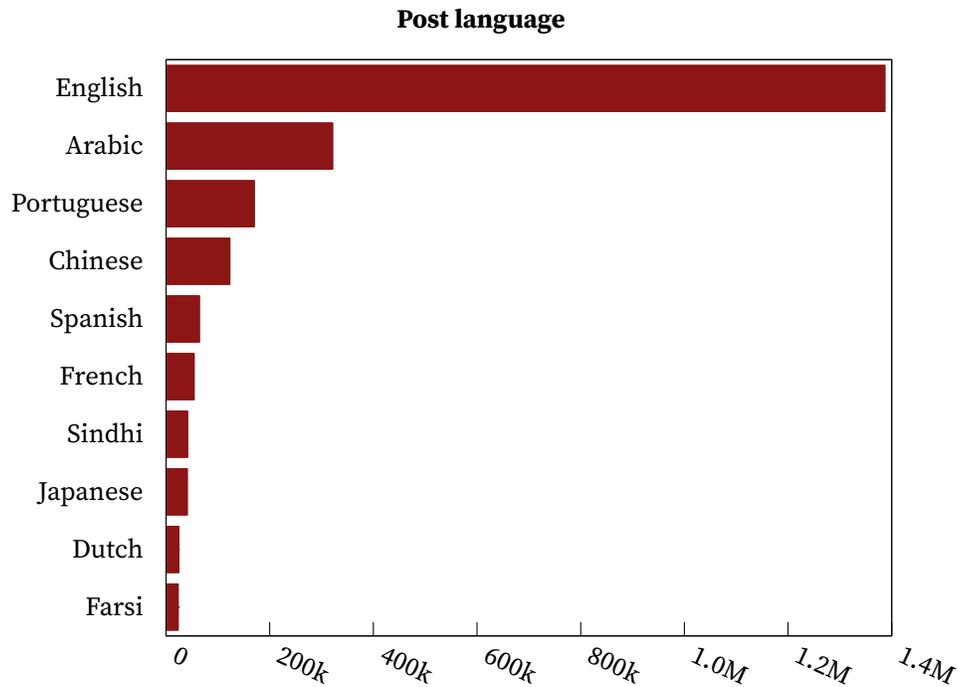


Figure 21: Top detected languages in post text, with high-false-positive categorizations removed. Some portion of Sindhi may be Arabic, and vice versa. Dutch false positives are 40–50%.

4.1.1 Most frequent post hashtags

In user posts, the most frequently used hashtags are primarily in reference to and in support of former US President Donald Trump, though a successful spam²⁴ campaign has led to #transrights and #transrightsarehumanrights appearing in the top 10. These are almost entirely the result of a single user posting the same comment repeatedly, which is apparently not detected by Gettr’s systems.

²⁴We use “spam” here not in the sense of commercial content, but rather repetitive and/or scripted high volume posting.

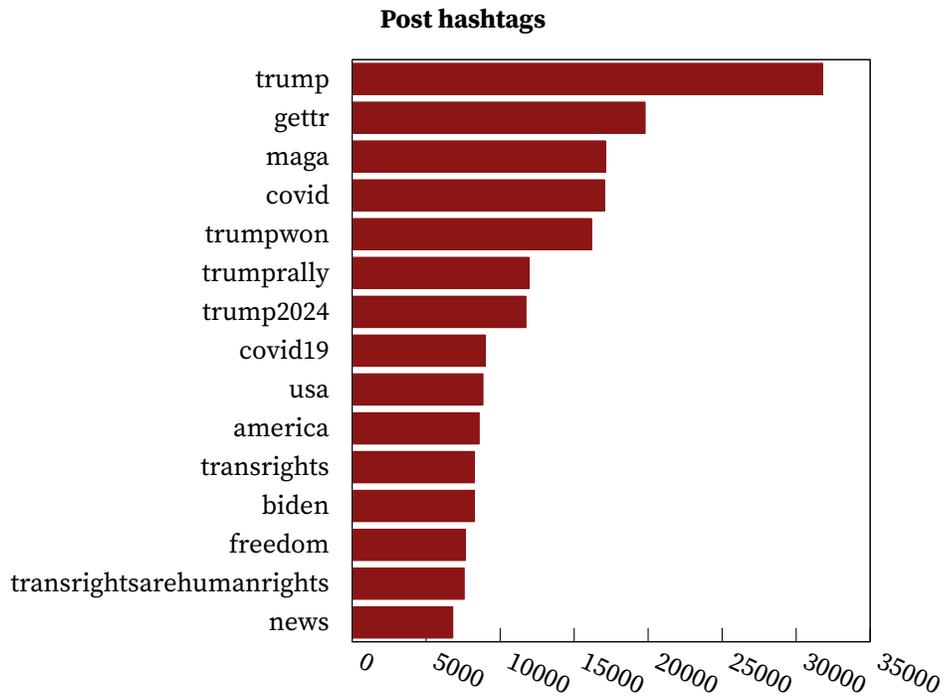


Figure 22: Most frequently used hashtags in user posts.

4.2 Emoji in posts

Emoji in post bodies themselves generally express a high proportion of “dunking,” religiosity, and patriotism for the US and Brazil. Interestingly, the loudly crying face is overwhelmingly used by Arabic-speaking users. English-speaking use of this emoji is for the most part humorous, rather than expressing sadness. Distribution of emoji in comments is largely the same, except with “rose” making the top 10 and the lack of “pointing down.”

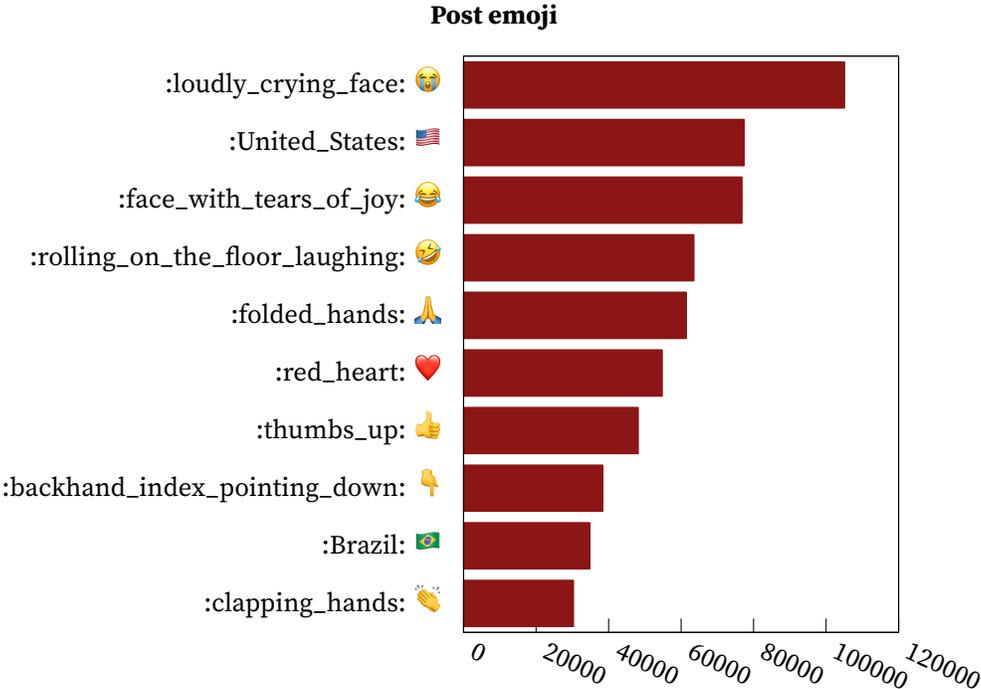


Figure 23: Top 10 most frequently used emoji in user posts.

4.3 Most frequently domains among posted links

Of posts including links, YouTube is the most highly represented, followed by Gateway Pundit, the Epoch Times and Guo’s GNews, along with a high number of Telegram groups. Largely unmoderated video sharing sites like Rumble, Bitchute and Odysee are also highly shared.

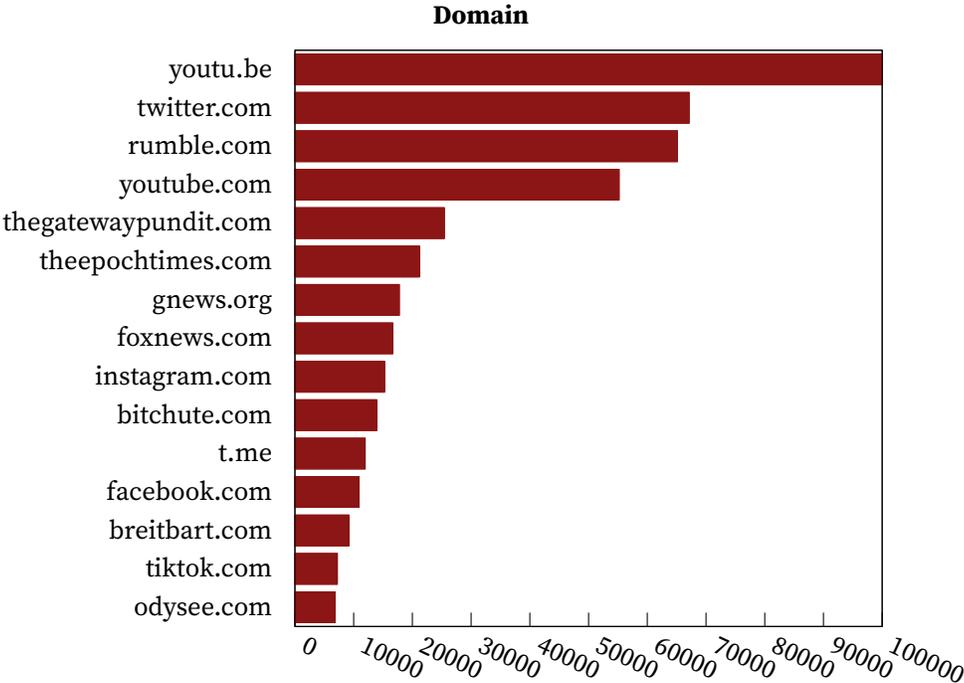


Figure 24: Domains of most frequently posted links.

4.4 Post frequency

The rate of new posts per minute has largely stabilized following a decline from the peak on July 5, indicating that Gettr is no longer growing (despite a small number of new user registrations every day; see Figure 9 on page 13). The pronounced vertical spike around July 24 is an aggressive spamming campaign; this can also be seen for comments in Figure 26 on the following page. The temporary dip in new posts on July 9 is presumably due to Gettr temporarily being down (or at least unable to accept new posts).

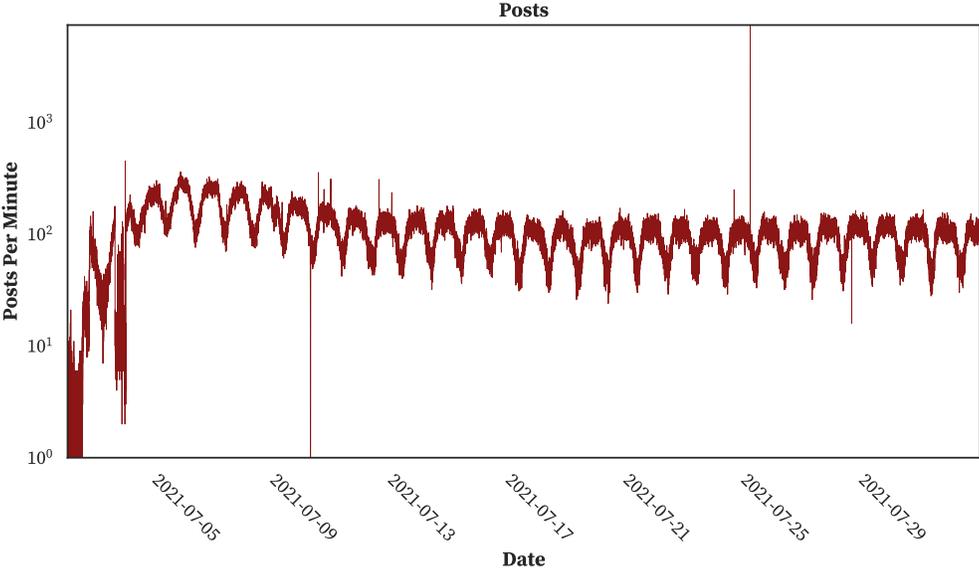


Figure 25: Posts per minute on Gettr during the month of July 2021 (log scale).

4.5 Comments

After the initial uptick in comment volume after launch, comments declined slowly before the unexplained jump in activity around July 13. As with posts, the large spikes show spam campaigns, as well as a downward spike presumably showing a brief outage.

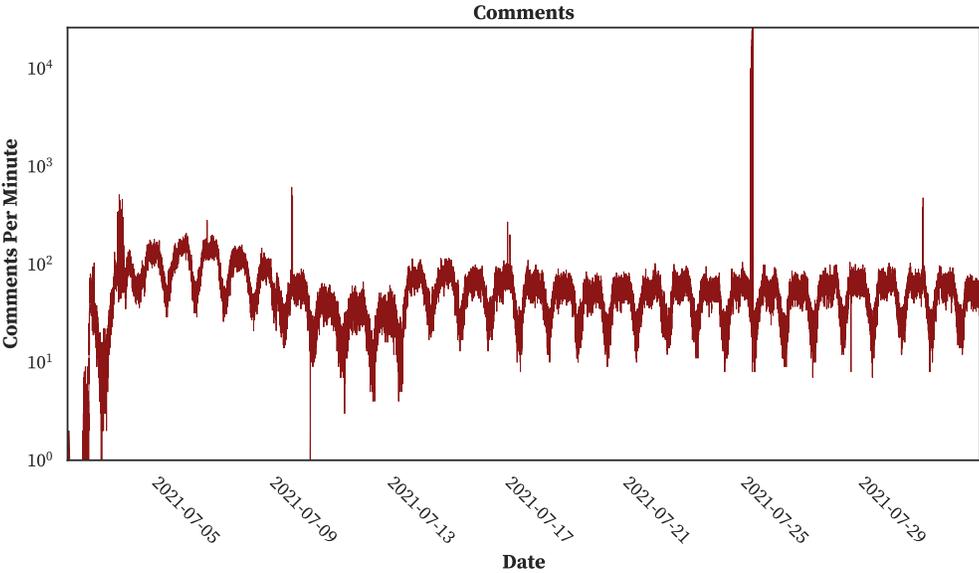


Figure 26: New comments per minute on Gettr during the month of July 2021 (log scale).

4.5.1 Hashtags

The most frequently used hashtags highlight the deficiencies in Gettr’s anti-spam measures. The most commonly used hashtag is #transrights, which as noted in Section 4.1.1 on page 22, is due to the efforts of a single party using an army of accounts with auto-generated usernames—a fairly trivial thing to detect. Other hashtags show a considerable amount of discourse being about Gettr itself, and other discourse referencing the Chinese Communist Party.

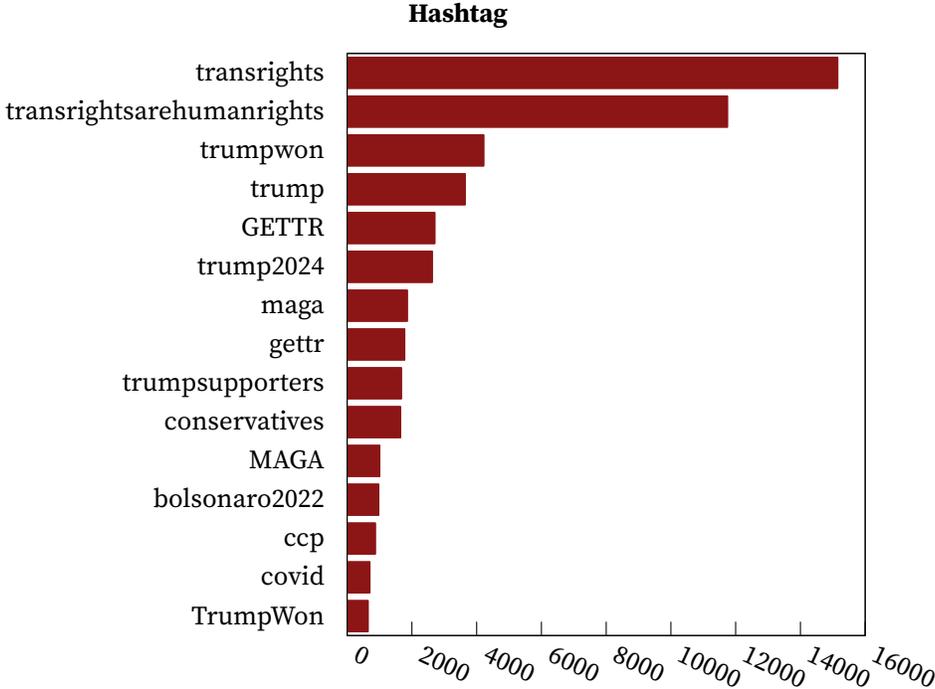


Figure 27: Most frequently used hashtags in user comments.

4.5.2 Replies and user tags

Users that garner the most replies are primarily level 4 influencers: Dinesh D’Souza, Miles Guo, Jason Miller, and Newsmax. Less obvious accounts are “Qanon211”—which joined shortly after the July 1 launch, and largely posts high volumes of generic right-wing third-party content regarding elections and the Biden administration—and “il Donaldo Trumpo,” an account posting similar content in faux-Spanish.

Highly tagged users in comments are a combination of popular accounts and personalities, along with relatively unknown but prolific “followback”/“IFBAP” users (see Figure 28 on the following page).



Figure 28: Two of the most highly tagged “followback” users.

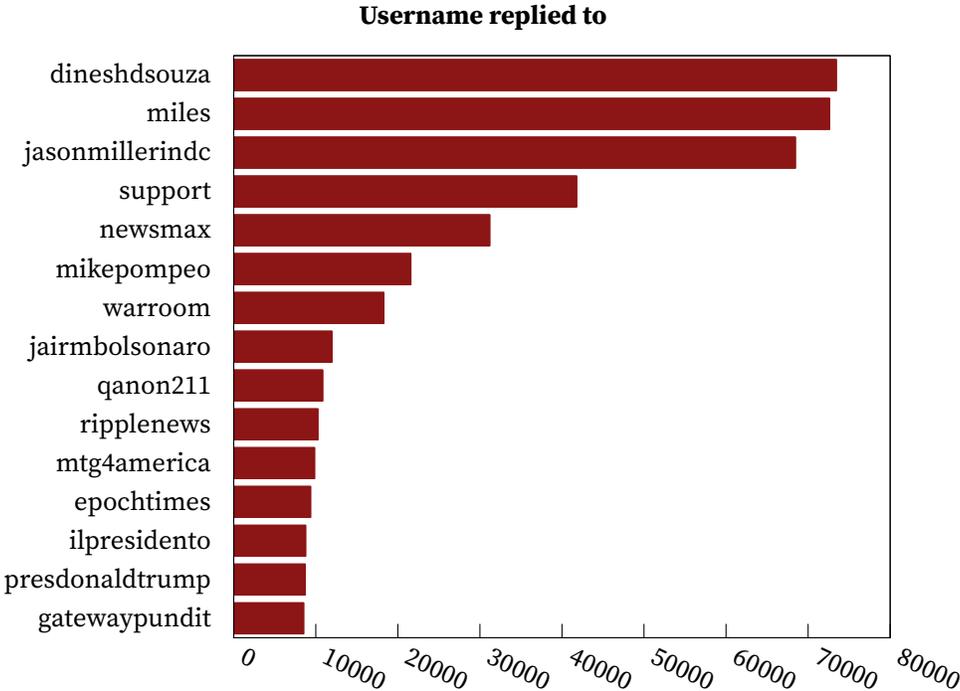


Figure 29: Users whose posts are most frequently replied to.

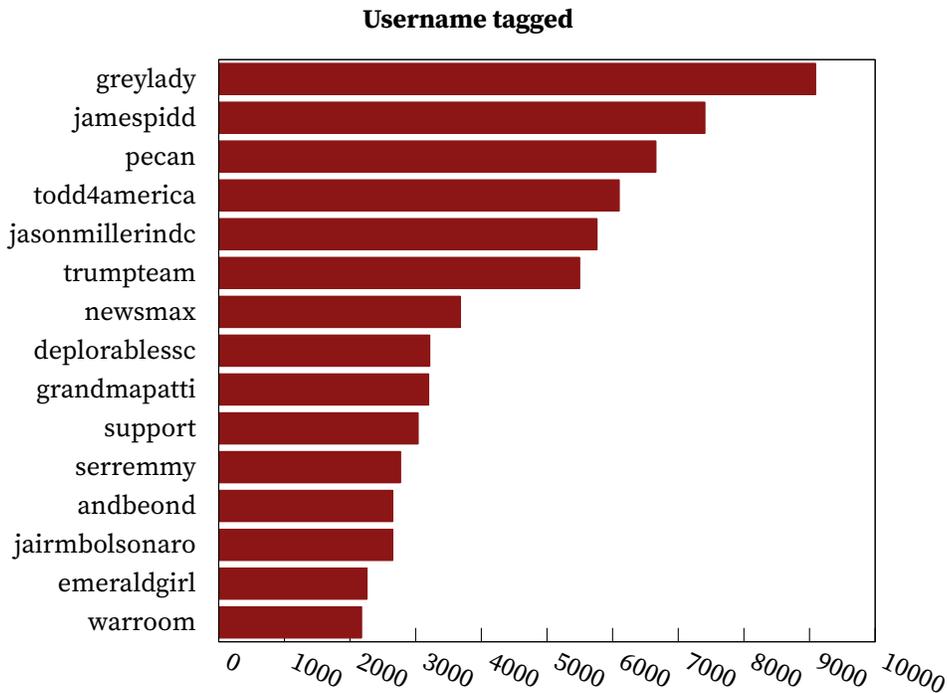


Figure 30: Users most frequently tagged in comments.

5 Sensitive and illegal content

Social media services generally use machine learning models to analyze uploaded image and video content to determine how to act on it—uploads can be rejected entirely, put behind a sensitive content filter or clickthrough, or, in severe cases, reported to law enforcement. As mentioned previously, Gettr does not appear to implement any kind of sensitive content detection—a survey of images using Google’s SafeSearch API indicates that 0.9% of posts with media and 1.8% of comments with media were classified as likely to contain violent or adult content (see Figures 31 and 32 on the next page), and as noted elsewhere, violent terrorist content has also surfaced on Gettr.²⁵

²⁵<https://www.politico.com/news/2021/08/02/trump-gettr-social-media-isis-502078>

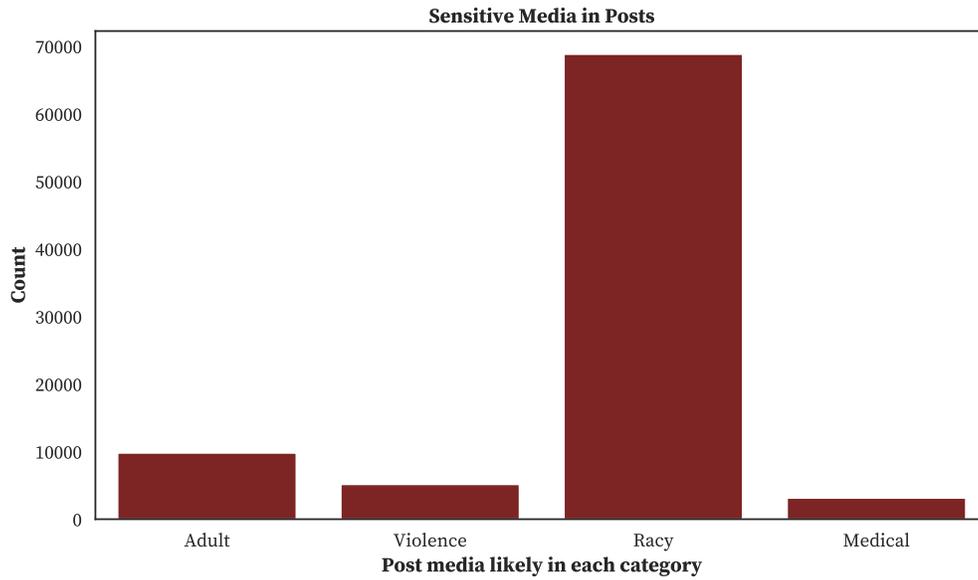


Figure 31: Total number of post images classified as “likely” or “very likely” to be sensitive.

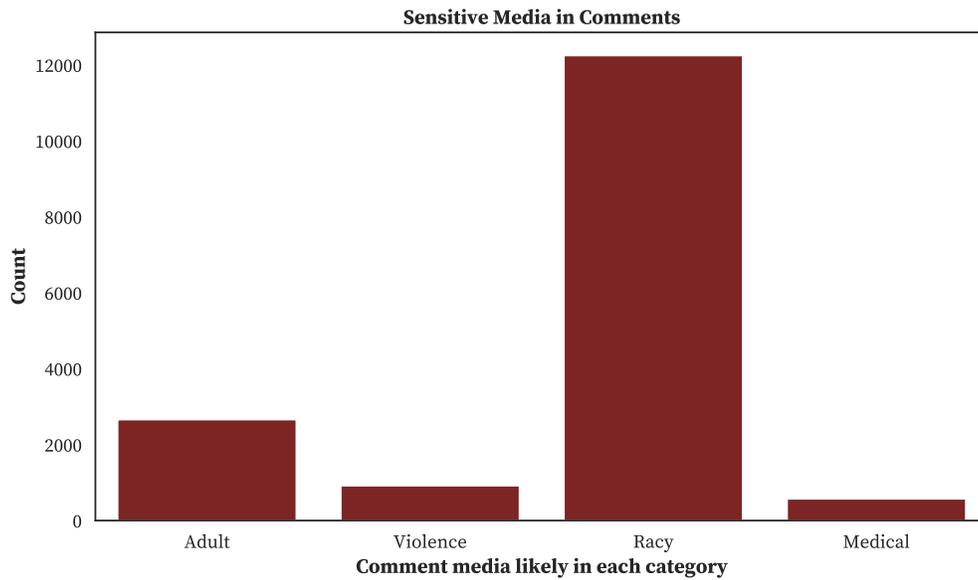


Figure 32: Total number of comment images classified as “likely” or “very likely” to be sensitive.

More importantly, Gettr also does not implement industry standard solutions for detection of Child Exploitation Imagery (CEI), most notably PhotoDNA.²⁶

Not only were we able to upload multiple images in the PhotoDNA test image set (a selection of benign images stored in the PhotoDNA database for testing purposes),

²⁶<https://en.wikipedia.org/wiki/PhotoDNA>

but a sampling of images from posts and comments on the site yielded 16 matches against production PhotoDNA datasets, which were subsequently reported by SIO to the National Center for Missing and Exploited Children (NCMEC).²⁷

Instead, Gettr appears to rely entirely on community reporting mechanisms to find sensitive content and illegal child-related imagery. An entirely community-based detection mechanism for sensitive and illegal content is extremely limited and prone to abuse, as such posts and comments may not be seen by users inclined to report them. Users may also not be aware of the reporting mechanisms themselves, or even what content qualifies as “child-related crime”—particularly given the fabricated child-related crime conspiracies that flourish on Gettr and similar platforms.

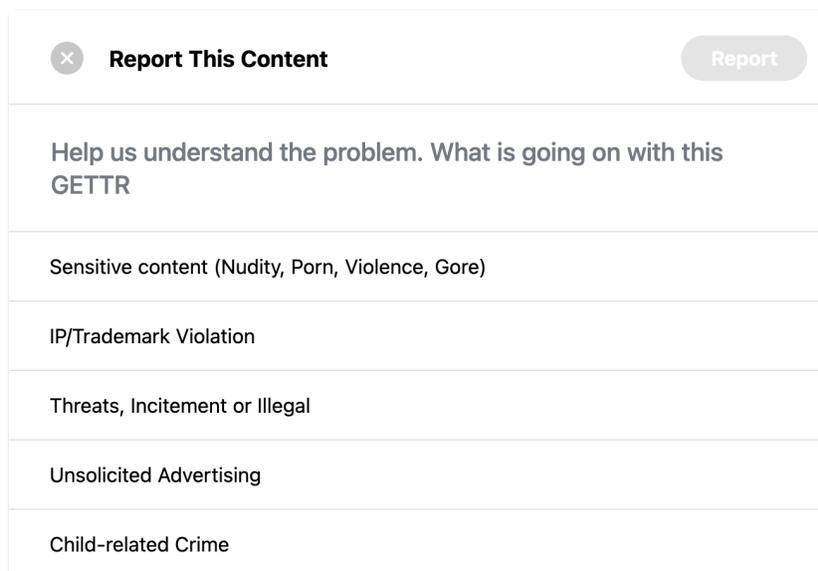


Figure 33: Gettr’s post and comment reporting dialog.

²⁷<https://www.missingkids.org/>

6 Conclusions

- Similar to other networks that cater to right wing and deplatformed communities, Gettr has a significant number of Portuguese and Arabic-speaking users (primarily Brazilian and Saudi); however, posting and commenting habits of those communities are not necessarily commensurate with their userbase.
- The user growth of Gettr appears to have been exaggerated, apparently having only reached 1.5 million at the beginning of August. Even with current numbers, spam accounts comprise a significant part of the userbase, and most accounts appear largely idle. In terms of posts and comments, growth has plateaued.
- The site seems to have been launched with little consideration for how to moderate content, and content policy does not appear cohesive or compliant with the terms of service. Porn remains on the platform while users were at launch prevented from using relatively mild expletives. We found no evidence of even cursory attempts at detecting spam and explicit imagery on the site.
- The site does not implement industry standard mechanisms to prevent the proliferation of known CEI. This is the bare minimum of content moderation systems, and no site which allows arbitrary uploads of media from the general public should be launched without it.

The Stanford Internet Observatory is a cross-disciplinary program of research, teaching and policy engagement for the study of abuse in current information technologies, with a focus on social media. The Stanford Internet Observatory was founded in 2019 to research the misuse of the internet to cause harm, formulate technical and policy responses, and teach the next generation how to avoid the mistakes of the past.

Stanford | Internet Observatory
Cyber Policy Center

